

Synthesis and Planning Under Environment Specifications

Antonio Di Stasio

distasio@diag.uniroma1.it

University of Rome "La Sapienza"

Reasoning Agents



ERC Advanced Grant

WhiteMech:

White-box Self Programming Mechanisms



SAPIENZA
UNIVERSITÀ DI ROMA



Given a specification φ over inputs I and outputs O , expressed in:


LTL (Pnueli 1977) or LTL_f (De Giacomo, Vardi 2013)

Syntax:

$$\varphi ::= a \mid \varphi \wedge \varphi \mid \neg \varphi \mid \bigcirc \varphi \mid \varphi \mathcal{U} \varphi \mid \Diamond \varphi \mid \Box \varphi$$

Semantic:

A trace $trace$ is an infinite (LTL) or finite (LTL_f) sequence over I and O . We write $trace \models \varphi$ to mean that τ satisfies φ .



Agent and Environment Strategies, and Traces

For an agent strategy $\sigma_{ag} : I^+ \rightarrow O$ and an environment strategy $\sigma_{env} : O^* \rightarrow I$, the trace

$$trace(\sigma_{ag}, \sigma_{env}) = (i_1 \cup o_1), (i_2 \cup o_2) \dots \in 2^{I \cup O}$$

denotes the unique trace induced by both σ_{ag} and σ_{env} .

Problem

Given an LTL/ LTL_f task *Goal* for the agent

Find agent strategy σ_{ag} such that $\forall \sigma_{env}. trace(\sigma_{ag}, \sigma_{env}) \models Goal$

LTL and LTL_f Synthesis

Algorithm for LTL synthesis

Given LTL formula φ

- 1: Compute corresponding NBA (Nondeterministic Buchi Aut.) (exponential)
- 2: Determinize NBA into DPA (Deterministic Parity Aut.) (exp in states, poly in priorities)
- 3: Synthesize winning strategy for Parity Game (poly in states, exp in priorities)

Algorithm for LTL_f synthesis

Given LTL_f formula φ

- 1: Compute corresponding NFA (Nondeterministic Finite Aut.) (exponential)
- 2: Determinize NFA to DFA (Deterministic Finite Aut.) (exponential)
- 3: Synthesize winning strategy for DFA game (linear)

Complexity

LTL and LTL_f synthesis are **2EXPTIME-complete**

Planning (or Synthesis with a model of the world)

Domain

- Planning consider the agent acting in a **(nondeterministic) domain**
- The domain is a **model of how the world** (i.e. the environment) works
- That is, it is a **specification of the possible environment strategies**

$$[[Dom]] = \{\sigma_{env} \mid \sigma_{env} \text{ compliant with } Dom\}$$

Planning in nondeterministic domains

Given an LTL_f task *Goal* for the agent, and a domain *Dom* modeling the environment

Find agent behavior σ_{ag} such that $\forall \sigma_{env} \in [[Dom]]. \text{trace}(\sigma_{ag}\sigma_{env}) \models Goal$

Specifying possible environment specifications in LTL/LTL_f

Environment specifications in LTL/LTL_f

Let Env be an LTL/LTL_f formula over I and O .

$$[[Env]] = \{\sigma_{env} \mid \forall \sigma_{ag}. trace(\sigma_{ag}, \sigma_{env}) \models Env\}$$

i.e. Env denotes all environment strategies that play according to the specification whatever is the agent strategy.

Synthesis under environment specifications in LTL/LTL_f

Given an LTL/LTL_f task $Task$ for the agent, and an LTL/LTL_f environment specification Env :

Find agent strategy σ_{ag} such that $\forall \sigma_{env} \in [[Env]]. trace(\sigma_{ag}, \sigma_{env}) \models Goal$

Environment specifications in LTL/LTLf

LTL/LTL_f for modeling the environment

1. Nondeterministic planning domains
2. Forms fairness ($\Box\Diamond\phi$) and stability ($\Diamond\Box\phi$) [ZhuDeGiacomoPuVardiAAAI20]
3. Safety properties. [DeGiacomoDiStasioPerelliZhuSubmitted]
4. General LTL environment specifications.[DeGiacomoDiStasioVardiZhuKR2020]
5. ...

Environment specifications in LTL/LTLf

LTL/LTL_f for modeling the environment

1. Nondeterministic planning domains
2. Forms fairness ($\Box\Diamond\phi$) and stability ($\Diamond\Box\phi$) [ZhuDeGiacomoPuVardiAAAI20]
3. Safety properties. [DeGiacomoDiStasioPerelliZhuSubmitted]
4. General LTL environment specifications. [DeGiacomoDiStasioVardiZhuKR2020]
5. ...

Consistent environment specifications

Is any LTL/LTL_f formula a valid environment specification? No, *Env* needs to be "consistent"!

$$[[Env]] \neq \emptyset \quad \text{i.e. } \exists \sigma_e. \forall \sigma_{ag}. \text{trace}(\sigma_{ag}, \sigma_e) \models Env$$

Synthesis Under Environment Specifications

Environment Specifications

Let Env be an LTL/LTL_f formula over $I \cup O$.

$$[[Env]] = \{\sigma_{env} \mid \sigma_{env} \text{ satisfies } Env \text{ whatever is the agent strategy}\}$$

Synthesis with environment specifications in LTL/LTL_f

Given an LTL/ LTL_f task $Goal$ for the agent, and an LTL/LTL_f environment specification Env :

Find agent strategy σ_{ag} such that $\forall \sigma_{env} \in [[Env]]. \text{trace}(\sigma_{ag}, \sigma_{env}) \models Goal$

Theorem [AminofDeGiacomoMuranoRubinICAPS2019]

To find agent strategy realizing $Goal$ under the environment specification Env , we can use standard LTL/LTL_f synthesis for

$$Env \rightarrow Goal$$

LTL_f Synthesis Under Safety Environment Specifications

Safety Properties

Definition

A safety property is a property which specifies that some (bad) behavior will never occur.

Examples:

- "always at most one process is in its critical section"
- "money can only be withdrawn once a correct PIN has been provided"

Important property

Any infinite trace violating the property has a finite prefix that is "bad";

- ... two processes are in the critical section ...

Usually: $\Box \neg \dots$

Consider a language $\mathcal{L} \subseteq \Sigma^\omega$.

Bad Prefix

A finite word $x \in \Sigma^*$ is a **bad prefix** for \mathcal{L} if and only if for all infinite words $y \in \Sigma^\omega$, we have $x \cdot y \notin \mathcal{L}$.

Safety Languages and Formulas

A language \mathcal{L} is a **safety language** iff every $w \notin \mathcal{L}$ has a bad prefix. A formula ϕ is a **safety formula** iff $\mathcal{L}(\phi) = \{w \in \Sigma^\omega \mid w \models \phi\}$ is a safety language.

Let $\mathcal{A} = (\Sigma, Q, q_0, \delta, F)$ be a NBA, we define its looping automaton $\mathcal{A}^{loop} = (\Sigma, Q \cup \{q_{sink}\}, q_0, \delta, Q)$, i.e., the automaton defined as \mathcal{A} in which every state has been made into an accepting state.

Theorem [KuperfmanVardiCAV99]

\mathcal{A} specifies a safety formula if and only if $\mathcal{L}(\mathcal{A}) = \mathcal{L}(\mathcal{A}^{loop})$.

Properties of looping automata:

1. Looping automata recognize safety languages.
2. Looping automata can be determinized by using **the standard subset construction**.

LTLf under Safety Environment Specifications

Problem

Solve the synthesis problem for

$$Env \rightarrow Goal$$

where *Env* is a safety environment specification expressed in LTL and *Goal* is an LTL_f formula.

Naive Solution

Translate to LTL and then do standard LTL synthesis for $Env \rightarrow Goal$.

... we can do better, since:

- *Env*: Safety

LTLf under Safety Environment Specifications

Problem

Solve the synthesis problem for

$$Env \rightarrow Goal$$

where *Env* is a safety environment specification expressed in LTL and *Goal* is an LTL_f formula.

Naive Solution

Translate to LTL and then do standard LTL synthesis for $Env \rightarrow Goal$.

... we can do better, since:

- *Env*: Safety
- *Goal*: LTL_f

1 ° Stage

1. Build the corresponding deterministic looping automaton $\mathcal{D}_{Env} = (\Sigma, S \cup \{s_{sink}\}, s_0, \delta, S)$ of Env .

1 ° Stage

1. Build the corresponding deterministic looping automaton $\mathcal{D}_{Env} = (\Sigma, S \cup \{s_{sink}\}, s_0, \delta, S)$ of Env .
2. Solve the safety game for the environment over \mathcal{D}_{Env} :

$$Safe_{env} = \nu Z. (S \cap Pre_{env}(Z))$$

$$Pre_{env}(Z) = \{s \mid \exists X \forall Y. \delta(s, X \cup Y) \in Z\}$$

1 ° Stage

1. Build the corresponding deterministic looping automaton $\mathcal{D}_{Env} = (\Sigma, S \cup \{s_{sink}\}, s_0, \delta, S)$ of Env .
2. Solve the safety game for the environment over \mathcal{D}_{Env} :

$$Safe_{env} = \forall Z. (S \cap Pre_{env}(Z))$$

$$Pre_{env}(Z) = \{s | \exists X \forall Y. \delta(s, X \cup Y) \in Z\}$$

3. Remove the winning region for the agent from \mathcal{D}_{Env} , say \mathcal{D}'_{Env} .

2 ° Stage

1. Build the corresponding DFA \mathcal{A}_{Goal}

2 ° Stage

1. Build the corresponding DFA \mathcal{A}_{Goal}
2. Do the cartesian product $\mathcal{G} = \mathcal{D}'_{Env} \times \mathcal{A}_{Goal}$.

2 ° Stage

1. Build the corresponding DFA \mathcal{A}_{Goal}
2. Do the cartesian product $\mathcal{G} = \mathcal{D}'_{Env} \times \mathcal{A}_{Goal}$.
3. Solve the reachability game for the agent over \mathcal{G} .

2 ° Stage

1. Build the corresponding DFA \mathcal{A}_{Goal}
2. Do the cartesian product $\mathcal{G} = \mathcal{D}'_{Env} \times \mathcal{A}_{Goal}$.
3. Solve the reachability game for the agent over \mathcal{G} .
4. Return the winning strategy for the agent if one exists.


LT_L^f Synthesis Under General LTL Environment Specifications

LTLf Synthesis Under LTL Environment Specifications

For example let the environment specifications be formed by $Env_1 \wedge Env_2$ where:

Env_1 is the LTL formula expressing the dynamics of the environment (as a planning domain):

- $\Box(\text{alive} \rightarrow \bigcirc(\text{wait} \rightarrow \text{alive}))$
- $\Box(\text{alive} \rightarrow \bigcirc(\text{shoot} \rightarrow (\text{alive} \vee \neg \text{alive})))$
- $\Box(\neg \text{alive} \rightarrow \bigcirc(\text{wait} \rightarrow \neg \text{alive}))$
- $\Box(\neg \text{alive} \rightarrow \bigcirc(\text{shoot} \rightarrow \neg \text{alive}))$
- $\Box((\text{wait} \wedge \text{shoot}) \wedge (\text{wait} \rightarrow \neg \text{shoot}))$




LTLf Synthesis Under LTL Environment Specifications

For example let the environment specifications be formed by $Env_1 \wedge Env_2$ where:

Env_1 is the LTL formula expressing the dynamics of the environment (as a planning domain):

- $\Box(\text{alive} \rightarrow \bigcirc(\text{wait} \rightarrow \text{alive}))$
- $\Box(\text{alive} \rightarrow \bigcirc(\text{shoot} \rightarrow (\text{alive} \vee \neg \text{alive})))$
- $\Box(\neg \text{alive} \rightarrow \bigcirc(\text{wait} \rightarrow \neg \text{alive}))$
- $\Box(\neg \text{alive} \rightarrow \bigcirc(\text{shoot} \rightarrow \neg \text{alive}))$
- $\Box((\text{wait} \wedge \text{shoot}) \wedge (\text{wait} \rightarrow \neg \text{shoot}))$



Env_2 is the LTL formula expressing some fairness over nondeterministic effects, e.g.,


$$\Box \Diamond \text{shoot} \rightarrow \Diamond \neg a$$

LTLf Synthesis Under LTL Environment Specifications

For example let the environment specifications be formed by $Env_1 \wedge Env_2$ where:

Env_1 is the LTL formula expressing the dynamics of the environment (as a planning domain):

- $\Box(alive \rightarrow \bigcirc(wait \rightarrow alive))$
- $\Box(alive \rightarrow \bigcirc(shoot \rightarrow (alive \vee \neg alive)))$
- $\Box(\neg alive \rightarrow \bigcirc(wait \rightarrow \neg alive))$
- $\Box(\neg alive \rightarrow \bigcirc(shoot \rightarrow \neg alive))$
- $\Box((wait \wedge shoot) \wedge (wait \rightarrow \neg shoot))$



Env_2 is the LTL formula expressing some fairness over nondeterministic effects, e.g.,

$$\Box \Diamond shoot \rightarrow \Diamond \neg a$$

Let $Goal$ be an LTL_f formula which expresses an agent task, e.g.,

$$\Diamond \neg a$$

LTLf Synthesis Under LTL Environment Specifications

Problem

Solve the synthesis problem for

$$Env_1 \wedge Env_2 \rightarrow Goal$$

LTLf Synthesis Under LTL Environment Specifications

Problem

Solve the synthesis problem for

$$Env_1 \wedge Env_2 \rightarrow Goal$$

Naive Solution

Translate to LTL and then do standard LTL synthesis for $Env_1 \wedge Env_2 \rightarrow Goal$.

LTLf Synthesis Under LTL Environment Specifications

Problem

Solve the synthesis problem for

$$Env_1 \wedge Env_2 \rightarrow Goal$$

Naive Solution

Translate to LTL and then do standard LTL synthesis for $Env_1 \wedge Env_2 \rightarrow Goal$.

... but we can exploit the simplicity of dealing with LTL_f given:

LTLf Synthesis Under LTL Environment Specifications

Problem

Solve the synthesis problem for

$$Env_1 \wedge Env_2 \rightarrow Goal$$

Naive Solution

Translate to LTL and then do standard LTL synthesis for $Env_1 \wedge Env_2 \rightarrow Goal$.

... but we can exploit the simplicity of dealing with LTL_f given:

- Env_1 : LTL

LTLf Synthesis Under LTL Environment Specifications

Problem

Solve the synthesis problem for

$$Env_1 \wedge Env_2 \rightarrow Goal$$

Naive Solution

Translate to LTL and then do standard LTL synthesis for $Env_1 \wedge Env_2 \rightarrow Goal$.

... but we can exploit the simplicity of dealing with LTL_f given:

- $Env_1 : LTL \rightarrow LTL_f$

LTLf Synthesis Under LTL Environment Specifications

Problem

Solve the synthesis problem for

$$Env_1 \wedge Env_2 \rightarrow Goal$$

Naive Solution

Translate to LTL and then do standard LTL synthesis for $Env_1 \wedge Env_2 \rightarrow Goal$.

... but we can exploit the simplicity of dealing with LTL_f given:

- $Env_1: LTL \rightarrow LTL_f$
- $Env_2: LTL$

LTLf Synthesis Under LTL Environment Specifications

Problem

Solve the synthesis problem for

$$Env_1 \wedge Env_2 \rightarrow Goal$$

Naive Solution

Translate to LTL and then do standard LTL synthesis for $Env_1 \wedge Env_2 \rightarrow Goal$.

... but we can exploit the simplicity of dealing with LTL_f given:

- $Env_1: LTL \rightarrow LTL_f$
- $Env_2: LTL$
- $Goal: LTL_f$

LTLf Synthesis Under LTL Environment Specifications

Separating LTL_f environment specifications

$$(Env_1 \wedge Env_2 \rightarrow Goal) \iff (Env_2 \rightarrow Env_1 \rightarrow Goal) \iff (Env_2 \rightarrow \neg Env_1 \vee Goal)$$

where $Goal' = \neg Env_1 \vee Goal$ is expressed in LTL_f and Env_2 in LTL.

LTLf Synthesis Under LTL Environment Specifications

Separating LTL_f environment specifications

$$(Env_1 \wedge Env_2 \rightarrow Goal) \iff (Env_2 \rightarrow Env_1 \rightarrow Goal) \iff (Env_2 \rightarrow \neg Env_1 \vee Goal)$$

where $Goal' = \neg Env_1 \vee Goal$ is expressed in LTL_f and Env_2 in LTL.

Problem

Solve the synthesis problem for

$$Env_2 \rightarrow Goal'$$

LTLf Synthesis Under LTL Environment Specifications

Separating LTL_f environment specifications

$$(Env_1 \wedge Env_2 \rightarrow Goal) \iff (Env_2 \rightarrow Env_1 \rightarrow Goal) \iff (Env_2 \rightarrow \neg Env_1 \vee Goal)$$

where $Goal' = \neg Env_1 \vee Goal$ is expressed in LTL_f and Env_2 in LTL.

Problem

Solve the synthesis problem for

$$Env_2 \rightarrow Goal'$$

How can we exploit that $Goal'$ is LTL_f?

LTLf Synthesis Under LTL Environment Specifications

Separating LTL_f environment specifications

$$(Env_1 \wedge Env_2 \rightarrow Goal) \iff (Env_2 \rightarrow Env_1 \rightarrow Goal) \iff (Env_2 \rightarrow \neg Env_1 \vee Goal)$$

where $Goal' = \neg Env_1 \vee Goal$ is expressed in LTL_f and Env_2 in LTL.

Problem

Solve the synthesis problem for

$$Env_2 \rightarrow Goal'$$


How can we exploit that $Goal'$ is LTL_f?

Two-stage technique!

Two-Stage Technique for Synthesis

1 ° Stage


1. Compute the corresponding DFA \mathcal{A} of $\neg Env_1 \vee Goal$.



Two-Stage Technique for Synthesis

1 ° Stage


1. Compute the corresponding DFA \mathcal{A} of $\neg Env_1 \vee Goal$.
2. Solve the reachability game for the agent over \mathcal{A} .



Two-Stage Technique for Synthesis

1 ° Stage


1. Compute the corresponding DFA \mathcal{A} of $\neg Env_1 \vee Goal$.
2. Solve the reachability game for the agent over \mathcal{A} .
3. Check whether the initial state is winning for the agent.



Two-Stage Technique for Synthesis

1 ° Stage


1. Compute the corresponding DFA \mathcal{A} of $\neg Env_1 \vee Goal$.
2. Solve the reachability game for the agent over \mathcal{A} .
3. Check whether the initial state is winning for the agent.
4. If the initial state is not winning go to Stage 2, otherwise return the agent winning strategy.



Two-Stage Technique for Synthesis

2 ° Stage



1. Remove from \mathcal{A} the agent winning set of Stage 1, say \mathcal{A}' .



Two-Stage Technique for Synthesis

2 ° Stage


1. Remove from \mathcal{A} the agent winning set of Stage 1, say \mathcal{A}' .
2. Compute the corresponding DPA \mathcal{B} of Env_2 .



Two-Stage Technique for Synthesis

2 ° Stage


1. Remove from \mathcal{A} the agent winning set of Stage 1, say \mathcal{A}' .
2. Compute the corresponding DPA \mathcal{B} of Env_2 .
3. Do the cartesian product between \mathcal{A}' and \mathcal{B} .



Two-Stage Technique for Synthesis

2 ° Stage


1. Remove from \mathcal{A} the agent winning set of Stage 1, say \mathcal{A}' .
2. Compute the corresponding DPA \mathcal{B} of Env_2 .
3. Do the cartesian product between \mathcal{A}' and \mathcal{B} .
4. Solve the parity game for the environment over $\mathcal{A}' \times \mathcal{B}$.



Two-Stage Technique for Synthesis

2 ° Stage


1. Remove from \mathcal{A} the agent winning set of Stage 1, say \mathcal{A}' .
2. Compute the corresponding DPA \mathcal{B} of Env_2 .
3. Do the cartesian product between \mathcal{A}' and \mathcal{B} .
4. Solve the parity game for the environment over $\mathcal{A}' \times \mathcal{B}$.
5. Check if the initial state is winning for the agent; if not return "Unrealizable".



Two-Stage Technique for Synthesis

2 ° Stage

1. Remove from \mathcal{A} the agent winning set of Stage 1, say \mathcal{A}' .
2. Compute the corresponding DPA \mathcal{B} of Env_2 .
3. Do the cartesian product between \mathcal{A}' and \mathcal{B} .
4. Solve the parity game for the environment over $\mathcal{A}' \times \mathcal{B}$.
5. Check if the initial state is winning for the agent; if not return "Unrealizable".
6. Return the agent winning strategy by combing the agent winning strategies in Stage 1 and 2.



We have

- implemented the two-stage technique in a new tool called **2SLS**, written in C++, that exploits CUDD package as library for the manipulation of Binary Decisions Diagrams (BDDs);
- compared **2SLS** to a direct reduction to LTL synthesis by employing the LTLf -to-LTL translator **SPOT** and **Strix** (Meyer, Sickert, and Luttenberger 2018) as the LTL synthesis solver;
- compared **2SLS** with FSyft and StSyft (Zhu et al. 2020) in special cases where environment specifications are LTL formulas of the form $\Box\Diamond a$ (fairness) and $\Diamond\Box a$ (stability), with a propositional.

Experiments on Fairness and Stability

- Given a counter game where the environment chooses whether to increment the counter or not and the agent can choose to grant the request or ignore it;
- The fairness specification is $\Box \Diamond \text{increment}$; the stability specification is $\Diamond \Box \text{increment}$;
- The goal is to get the counter having all bits set to 1.




Figure 1: LTL_f synthesis under fairness specifications.




Figure 2: LTL_f synthesis under stability specifications.

Experiments of General LTL Environment Specifications

- Given *Goal* as a conjunction of increasing size of random LTL_f formulas of the form $\Box(p_j \rightarrow \Diamond q_j)$ with p_j and q_j propositions under the control of the environment and the agent, respectively;
- Env* is a conjunction of formulas of the form $(\Box \Diamond p_i \vee \Diamond \Box q_i)$, where we start with one conjunct and introduce a new conjunct every 10 conjuncts in *Goal*.




Figure 3: LTL_f synthesis under general LTL environment specifications.

Develop LTL_f synthesis under safety environment specifications where you can use:

- Spot (spot.lrde.epita.fr) for building the looping deterministic automaton
- LTLf2DFA (ltlf2dfa.diag.uniroma1.it) for constructing the DFA.