

LTL_f -based Trace Alignment: a Planning-based Approach

Fabio Patrizi¹

Sapienza University of Rome, Italy

patrizi@diag.uniroma1.it



Agents-VIC
May 7, 2021

¹Jointly with G. De Giacomo, F.M. Maggi, A. Marrella, A. Murano, G. Perelli

Process Mining: analysis of (business) processes starting from *event logs* [van der Aalst, 2016]

Event Log



Events model process activities:

a =register request, b =examine thoroughly

c =examine casually, d = check ticket

e =decide, f =reinitiate request,

g =pay compensation, h =reject request

Traces: event sequences modeling a process run

1. a, b, d, e, h

2. a, d, c, e, h

3. $a, c, d, e, f, b, d, e, g$

4. a, d, b, e, h

...

Event Log: set of traces

case id	event id	properties		
		timestamp	activity	resource
1	35654423	30-12-2010:11.02	register request	Pete
	35654424	31-12-2010:10.06	examine thoroughly	Sue
	35654425	05-01-2011:15.12	check ticket	Mike
	35654426	06-01-2011:11.18	decide	Sara
	35654427	07-01-2011:14.24	reject request	Pete
2	35654483	30-12-2010:11.32	register request	Mike
	35654485	30-12-2010:12.12	check ticket	Mike
	35654487	30-12-2010:14.16	examine casually	Pete
	35654488	05-01-2011:11.22	decide	Sara
	35654489	08-01-2011:12.05	pay compensation	Ellen
3	35654521	30-12-2010:14.32	register request	Pete
	35654522	30-12-2010:15.06	examine casually	Mike
	35654524	30-12-2010:16.34	check ticket	Ellen
	35654525	06-01-2011:09.18	decide	Sara
	35654526	06-01-2011:12.18	reinitiate request	Sara
	35654527	06-01-2011:13.06	examine thoroughly	Sean
	35654530	08-01-2011:11.43	check ticket	Pete
	35654531	09-01-2011:09.55	decide	Sara
	35654533	15-01-2011:10.45	pay compensation	Ellen
4	35654641	06-01-2011:15.02	register request	Pete
	35654643	07-01-2011:12.06	check ticket	Mike
	35654644	08-01-2011:14.43	examine thoroughly	Sean
	35654645	09-01-2011:12.02	decide	Sara
	35654647	12-01-2011:15.44	reject request	Ellen

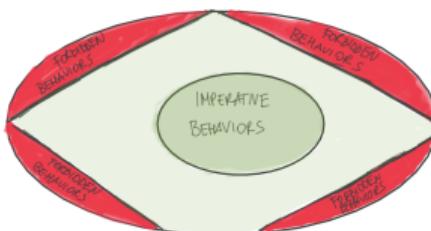
Typical problems in Process Mining:

- Process Discovery: discover the underlying process that produced the log [Cook and Wolf, 1995]
- Conformance Checking: check whether a log conforms to a given process model [Rozinat and van der Aalst, 2008]
- Monitoring: monitor the current trace wrt a given process model [Ly et al., 2013]

Declarative Process Model

Imperative Process Model

- Prescriptive specification
- Process fully specified
- Common Models:
 - Petri Net
 - Finite-state automaton
 - BPMN
 - ...



- Descriptive specification
- Behaviors subject to constraints
- Behaviors allowed unless violating
- Constraints expressed in:
 - DECLARE [Pescic et al., 2007]
 - LTL_f/LDL_f [De Giacomo and Vardi, 2013]

Declarative Process Mining (this talk): Event Log + Declarative Process Model

Trace Alignment is a form of Conformance Checking

- Traces often *dirty* (involve human activities)
 - *spurious* or *missing* events
 - in general, *discrepancies* may be present
- Process model not enforced

Trace Alignment

Trace Alignment is the problem of *cleaning* and *repairing* dirty traces to make them compliant with the underlying process model [van der Aalst, 2016]

Solving the problem allows for:

- Identifying trace errors
- Taking corrective actions
- Measuring deviation wrt process model

Trace Alignment [Bose and van der Aalst, 2010] (declarative variant)

Given:

- A trace ρ over a finite event alphabet $\Sigma = \{\sigma_1, \dots, \sigma_n\}$
- A constraint φ (expressed in some formal language, see below)
- A cost function $cost$ associating non-negative costs to *additions* and *deletions* of each event

Find trace ρ' over Σ s.t.:

- ρ' satisfies φ , written $\rho \models \varphi$
- Cost $cost(\rho, \rho')$ of turning ρ into ρ' is minimal (wrt the cost of changes made to ρ)

Example

- $\rho = a \ a \ c \ b \ a \ c \ a$
- φ = “whenever a occurs, b eventually occurs after a ”
- all changes cost 1
- ρ does not satisfy φ (last a has no matching b)
 - ① Remove all a 's: $\rho_1 = a \ a \ c \ b \ a \ c \ a = c \ b \ c$, $cost(\rho, \rho_1) = 4$
 - ② Add b after every a : $\rho_2 = a \ b \ a \ b \ c \ b \ a \ b \ c \ a \ b$, $cost(\rho, \rho_2) = 4$
 - ③ Add b at end only: $\rho_3 = a \ a \ c \ b \ a \ c \ a \ b$, $cost(\rho, \rho_3) = 1$ (solution!)

Observations:

- if φ satisfiable, repaired trace always exists
- we can think of $cost(\rho, \rho')$ as the distance of ρ from φ (if $\rho' \models \varphi$)

Some uses of trace alignment in BP:

- Check whether expected process is correctly executed and quantify deviation (costs)
- Identify and correct errors in process execution
- Identify common patterns among log traces

But also in AI:

- Use traces to model agent behaviors (e.g., executed actions)
- Use trace alignment to:
 - compare observed behavior with expected/desired one
 - quantify discrepancies and repair
 - identify behavioral patterns
 - ...

- ➊ A technique for trace alignment [De Giacomo et al., 2017]:
 - automata-based approach
 - solved with cost-optimal planning
- ➋ Experimental results
- ➌ (Extension to metric temporal constraints)

Constraint Specification

DECLARE



Standard language in Declarative PM is DECLARE [Pesic et al., 2007]

Defined through *templates*:

- *Existence(a)*: activity a is executed at least once
 $b\ a\ c\ a, b\ c\ c$
- *Response(a, b)*: if a occurs then b eventually occurs after a
 $a\ c\ a, a\ c\ a\ b, c\ d\ c$
- *Choice(a, b)*: a or b (possibly both) eventually occur
 $a\ c\ a, a\ c\ b, c\ c\ b, d\ c\ c$
- *ChainResponse(a, b)*: whenever a occurs, b immediately follows
 $a\ c\ c\ b, d\ c\ c, d\ a\ b)$
- ...

Constraint Specification

LTL_f

DECLARE is a fragment of LTL_f

LTL_f (linear-time temporal logic over finite traces [De Giacomo and Vardi, 2013])

$\varphi := \text{true} \mid a \mid \neg\varphi \mid \varphi \wedge \varphi \mid \mathbf{X}\varphi \mid \varphi \mathbf{U} \varphi$

- $\mathbf{X}\varphi$ (“next φ ”): next event satisfies φ
- $\varphi_1 \mathbf{U} \varphi_2$ (“ φ_1 until φ_2 ”): an event satisfying φ_2 eventually occurs (possibly now) and all events until then satisfy φ_1

Standard abbreviations (Boolean as usual):

- $\mathbf{F}\varphi \doteq \text{true} \mathbf{U} \varphi$ (“eventually φ ”): φ eventually holds
- $\mathbf{G}\varphi \doteq \neg \mathbf{F} \neg \varphi$ (“always φ ”): φ always holds

Constraint Specification

DECLARE and LTL_f

Some examples

DECLARE template	LTL_f translation
<i>Existence(a)</i>	$\mathbf{F}a$
<i>Response(a, b)</i>	$\mathbf{G}(a \rightarrow \mathbf{F}b)$
<i>Choice(a, b)</i>	$\mathbf{F}a \vee \mathbf{F}b$
<i>ChainResponse(a, b)</i>	$\mathbf{G}(a \rightarrow \mathbf{X}b)$

We use full LTL_f (in fact, could go even beyond and use LDL_f)

Example

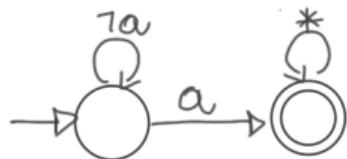
- $\rho = a \ a \ c \ b \ a \ c \ a$
- $\varphi = \mathbf{G}(a \rightarrow \mathbf{F}b)$ (“whenever a occurs, b eventually occurs after a ”)
- ...

Theorem ([De Giacomo and Vardi, 2013])

Every LTL_f formula φ has a corresponding (exponential) NFA \mathcal{A}_φ s.t.

for every trace ρ : $\rho \models \varphi \Leftrightarrow \mathcal{A}_\varphi$ accepts ρ

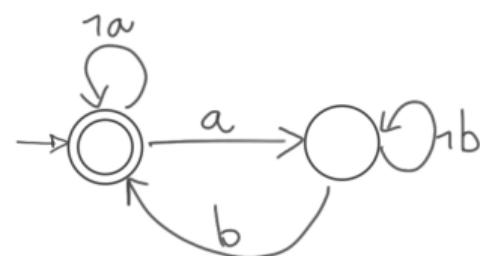
$$\varphi = \mathbf{F}a$$



$$\varphi = \mathbf{G}a$$



$$\varphi = \mathbf{G}(a \rightarrow \mathbf{F}b)$$



Automata-based Solution

Idea

Given:

- trace ρ
- constraint φ (one, w.l.o.g.)

Define:

- *Augmented trace automaton* \mathcal{T}^+ : accepts all modifications of ρ
- *Augmented constraint automaton* \mathcal{A}^+ :
accepts all traces that satisfy φ **plus** all modifications of ρ that satisfy φ

Find minimal-cost ρ' s.t.:

- ρ' is accepted by \mathcal{T}^+
- ρ' is accepted by \mathcal{A}^+

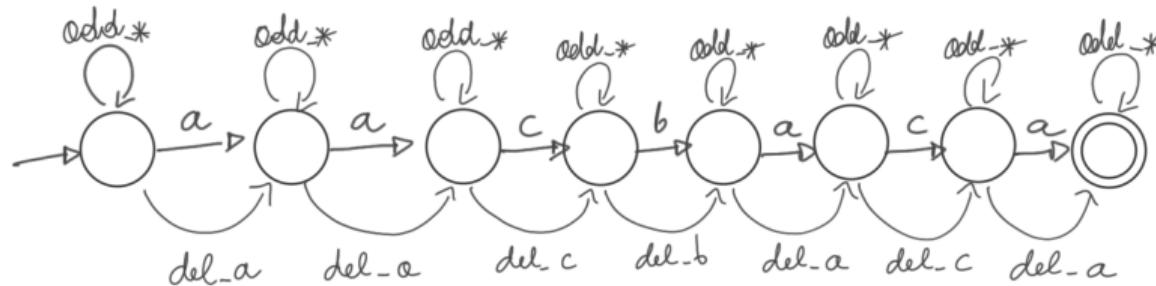
Augmented trace automaton \mathcal{T}^+



SAPIENZA
UNIVERSITÀ DI ROMA

Example:

- $\rho = a \ a \ c \ b \ a \ c \ a$



Accepts all modifications of ρ , with changes **marked**, e.g.:

- ① Remove all a's: **del_a del_a c b del_a c del_a**
- ② Add b after every a: **a add_b a add_b c b a add_b c a add_b**
- ③ Add b at end only: **a a c b a c a add_b**

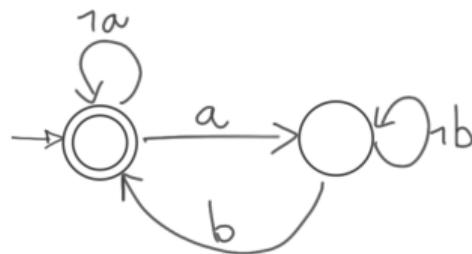
Augmented constraint automaton \mathcal{A}^+



Example:

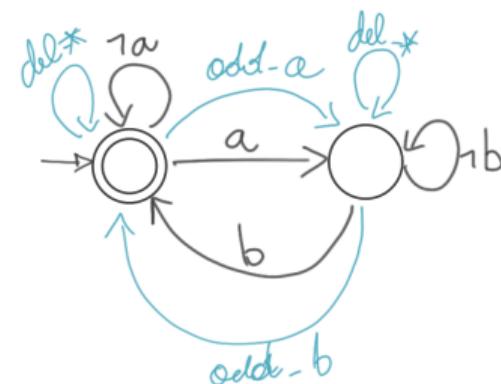
- $\varphi = \mathbf{G}(a \rightarrow \mathbf{F}b)$ ("whenever a occurs, b eventually occurs after a ")

Plain Constraint Automaton



\Rightarrow

Augmented Constraint automaton



Observe: accepts also traces that are not modifications of ρ

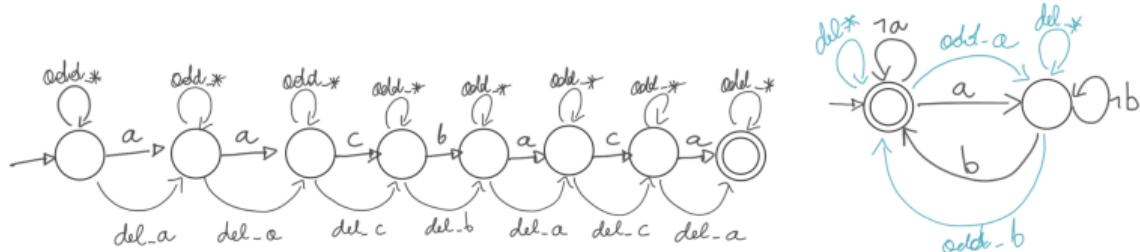
Problem

Find **miminal-cost** ρ' s.t.:

- ρ' is accepted by \mathcal{T}^+ , i.e., changes wrt ρ are marked by adds and dels
- ρ' is accepted by \mathcal{A}^+ , i.e., ρ' “repairs” ρ

Corresponds to finding minimal-cost accepting path ρ' on product automaton $\mathcal{T}^+ \times \mathcal{A}^+$

- $\rho = a\ a\ c\ b\ a\ c\ a$, $\varphi = \mathbf{G}(a \rightarrow \mathbf{F}b)$
- $\rho' = a\ a\ c\ b\ a\ c\ a$ **add_b**



Theorem

Trace alignment can be solved in time polynomial wrt $|\rho| \times 2^{|\varphi|}$

- *Searching min-cost path in $\mathcal{T}^+ \times \mathcal{A}^+$: $\mathcal{O}(N \log N)$, with $N = |\mathcal{T}^+ \times \mathcal{A}^+| = |\mathcal{T}^+| \times |\mathcal{A}^+|$*
- $|\mathcal{T}^+| = \mathcal{O}(|\rho|)$, *time*: $\mathcal{O}(|\rho|)$
- $|\mathcal{A}_\varphi| = \mathcal{O}(2^{|\varphi|})$, *time*: $\mathcal{O}(2^{|\varphi|})$
- $|\mathcal{A}^+|$: $\mathcal{O}(|\mathcal{A}_\varphi|)$, *time*: $\mathcal{O}(|\mathcal{A}_\varphi|)$

Use planning to search for minimal-cost ρ'

- Domain:
 - Models product automaton $\mathcal{T}^+ \times \mathcal{A}^+$
 - *sync* actions with null cost model events
 - *add* and *del* actions with positive costs model changes to input trace
 - Dynamics model *synchronous* execution of all augmented automata (trace+constraints)
- Problem:
 - Initial state: all automata in their starting state
 - Goal: all automata in a final state
- Solution:
 - Minimal-cost goal-reaching sequence of actions

Use planning to search for minimal-cost ρ'

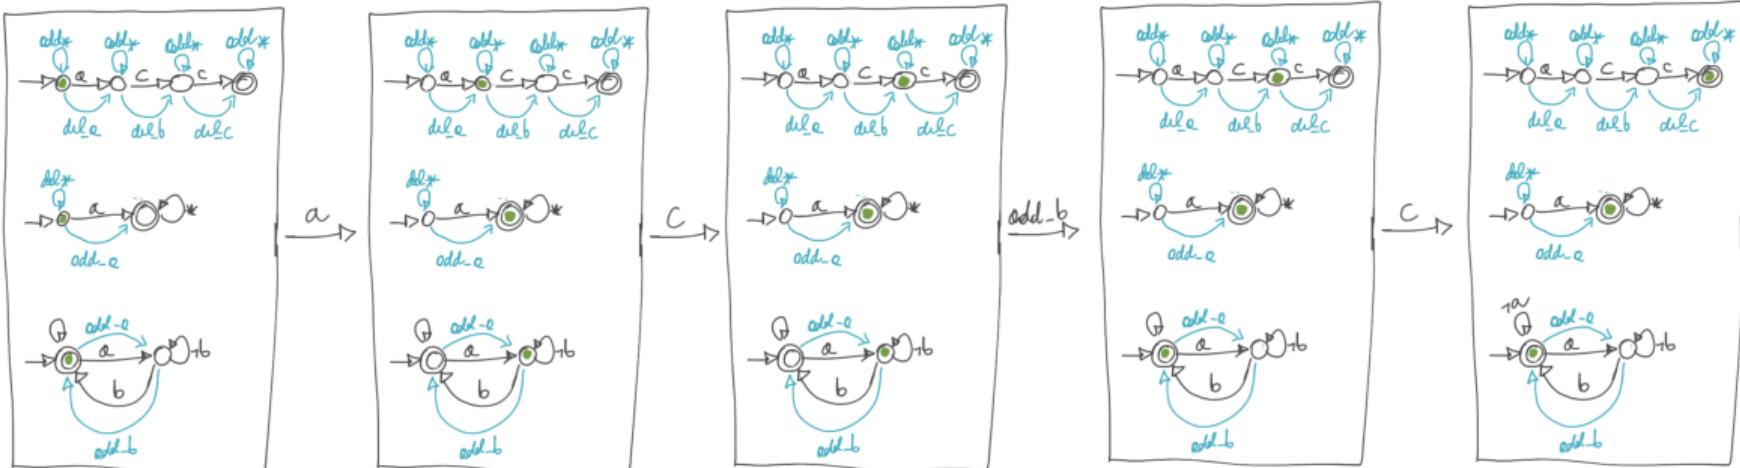
- Actions with costs model events and changes to ρ
- Adds and dels have (possibly different) positive cost
- Augmented automata (trace and constraints) modeled by domain
- Solution: sequence of actions that takes all automata to accepting state

TASK: find minimal-cost plan

Trace alignment as minimal-cost planning



- $\rho = a \ c \ c$
- $\varphi_1 = \mathbf{F}a$
- $\varphi_2 = \mathbf{G}(a \rightarrow \mathbf{F}b)$
- $\rho' = a \ c \ \text{add_}b \ c$



Solving trace alignment with planning



The problem can be encoded in PDDL 2.1 [Fox and Long, 2003]

```
(define (domain alignment)
  (:requirements :typing :disjunctive-preconditions :conditional-effects
  :universal-preconditions :action-costs)
  (:types trace_state automaton_state - state activity)
  (:predicates
    (trace ?t1 - trace_state ?e - activity ?t2 - trace_state)
    (cur_state ?s - state)
    (automaton ?s1 - automaton_state ?e - activity ?s2 - automaton_state)
    (final_state ?s - state)
  )
  (:functions total-cost)

  (:action sync
    :parameters (?t1 - trace_state ?e - activity ?t2 - trace_state)
    :precondition (and (cur_state ?t1)(trace ?t1 ?e ?t2))
    :effect(and
      (not (cur_state ?t1))
      (cur_state ?t2)
      (forall (?s1 ?s2 - automaton_state)
        (when (and (cur_state ?s1) (automaton ?s1 ?e ?s2))
          (and (not (cur_state ?s1))(cur_state ?s2)))
        )
      )
    )
  )

  (:action add
    :parameters (?e - activity)
    :effect (and
      (increase (total-cost) 1)
      (forall (?s1 ?s2 - automaton_state)
        (when (and (cur_state ?s1) (automaton ?s1 ?e ?s2))
          (and (not (cur_state ?s1))(cur_state ?s2)))
        )
      )
    )
  )

  (:action del
    :parameters (?t1 - trace_state ?e - activity ?t2 - trace_state)
    :precondition (and (cur_state ?t1)(trace ?t1 ?e ?t2))
    :effect(and
      (increase (total-cost) 1)
      (not (cur_state ?t1)) (cur_state ?t2)
    )
  )
)
```

Various encodings are possible, huge impact on performance

Experiments



- Tested (cost-optimal) planners:
 - Fast-Downward [Helmert, 2006]
 - SymBA* [Torralba et al., 2016]
 - (Assumed unitary costs for adds and dels)
- Compared against
 - Ad-hoc SOA approach implemented in ProM
[de Leoni and van der Aalst, 2013, de Leoni et al., 2015]
- Dataset
 - Real-life logs (loan application in Dutch financial institute), 16 constraints
 - Synthetic logs with 10, 15 and 20 constraints

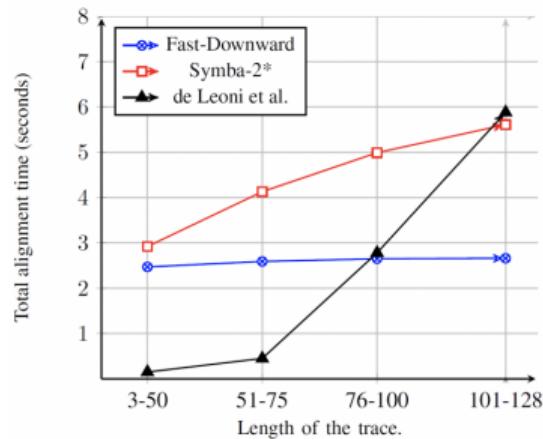
Results

planning vs ad-hoc on real-life logs



SAPIENZA
UNIVERSITÀ DI ROMA

trace length	no. traces	Fast-Downward	SymbA-2*	de Leoni et al.	Alignment Cost
<i>Real-life log</i>		16 constraints			
3-50	607	2.47	2.92	0.15	0.63
51-75	38	2.59	4.13	0.45	1.02
76-100	5	2.65	4.99	2.78	2.4
101-128	4	2.66	5.61	5.88	2.5

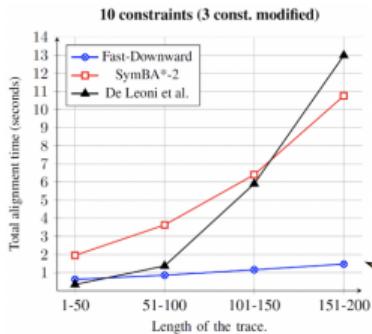


Results

planning vs ad-hoc on synthetic logs

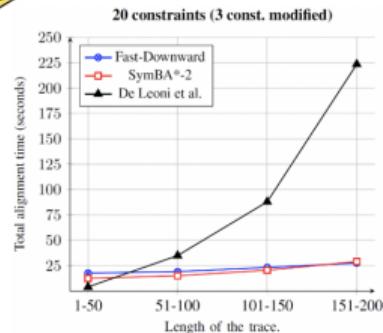
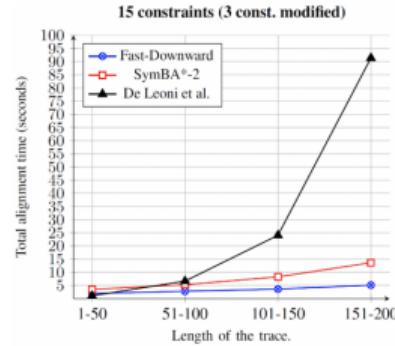


Trace length	Fast-Downward	SymbA-2*	de Leoni et al.	Alignment Cost
3 const. modified				
10 constraints				
1-50	0.62	1.95	0.34	1.77
51-100	0.85	3.63	1.37	2.11
101-150	1.15	6.4	5.9	3.03
151-200	1.46	10.75	12.98	3.79
15 constraints				
1-50	1.97	3.49	1.08	1.71
51-100	2.79	5.3	6.64	2.23
101-150	3.61	8.26	24.05	3.07
151-200	5.12	13.63	91.39	4.2
20 constraints				
1-50	17.63	12.42	3.99	1.87
51-100	19.05	15.02	34.91	2.61
101-150	23.23	20.45	87.89	3.35
151-200	27.49	28.97	223.47	4.2



Our approach **outperforms ProM** by **several orders of magnitude** when the model becomes larger and the noise increases.

The ad-hoc approach implemented in ProM is faster only for short traces with little noise.



- ➊ Automata-based approach implemented with planning technology
(available at: tinyurl.com/glymp9d)
- ➋ Outperforms ad-hoc solutions
 - More scalable
 - More flexible (other planners can be plugged)
- ➌ PDDL encoding has huge impact, usually little preprocessing required –great benefit

- Extending the planning-based approach to LDL_f [De Giacomo and Vardi, 2013], i.e., LTL_f with RE
(as expressive as FSA)
- Testing new encodings
- Decidability of the problem for MTL_f [Koymans, 1990], the time-metric extension of LTL_f
(see below)

Summing up:

- Problem originated in Declarative Process Mining
- Practical and Speculative Relevance also to AI (Agent Behaviors)
- Two variants considered: untimed and timed
- AI-planning technique provides best results in untimed setting
- Timed setting solvable but practicality to be assessed

Future work:

- Extend untimed version to LDL_f and new PDDL encodings (ongoing)
- Study well-behaved fragments of MTL_f or effects of constraints on repairs (e.g., fixed number of adds within two events)
- Implement solution approaches for timed variant (planning-based?)
- Consider data-aware setting (where activities carry a payload) –particularly relevant to KR

Thank you!



SAPIENZA
UNIVERSITÀ DI ROMA

Questions?

References I



Bose, R. P. J. C. and van der Aalst, W. M. P. (2010).

Trace alignment in process mining: Opportunities for process diagnostics.

In *Proc. of BPM '10*.



Cook, J. E. and Wolf, A. L. (1995).

Automating process discovery through event-data analysis.

In *Proc. of ICSE'95*.



De Giacomo, G., Maggi, F. M., Marrella, A., and Patrizi, F. (2017).

On the disruptive effectiveness of automated planning for LTL_f-based trace alignment.

In *Proc. of AAAI'17*.



De Giacomo, G. and Vardi, M. Y. (2013).

Linear temporal logic and linear dynamic logic on finite traces.

In *Proc. of IJCAI'13*.



de Leoni, M., Maggi, F. M., and van der Aalst, W. M. P. (2015).

An alignment-based framework to check the conformance of declarative process models and to preprocess event-log data.

Inf. Syst., 47:258–277.



de Leoni, M. and van der Aalst, W. M. P. (2013).

Aligning event logs and process models for multi-perspective conformance checking: An approach based on integer linear programming.

In *Proc. of BPM '13*.



Finkel, A. and Schnoebelen, P. (2001).

Well-structured transition systems everywhere!

Theor. Comput. Sci., 256(1-2):63–92.

References II



Fox, M. and Long, D. (2003).
PDDL2.1: An Extension to PDDL for Expressing Temporal Planning Domains.
J. Artif. Intell. Res. (JAIR), 20:61–124.



Helmer, M. (2006).
The fast downward planning system.
J. Artif. Intell. Res., 26:191–246.



Henriksen, J. G., Jensen, J. L., Jørgensen, M. E., Klarlund, N., Paige, R., Rauhe, T., and Sandholm, A. (1995).
Mona: Monadic second-order logic in practice.
In *Proc. of TACAS '95*.



Koymans, R. (1990).
Specifying real-time properties with metric temporal logic.
Real Time Syst., 2(4):255–299.



Ly, L. T., Maggi, F. M., Montali, M., Rinderle-Ma, S., and van der Aalst, W. M. P. (2013).
A framework for the systematic comparison and evaluation of compliance monitoring approaches.
In *Proc. of EDOC'13*.



Ouaknine, J. and Worrell, J. (2007).
On the Decidability and Complexity of Metric Temporal Logic over Finite Words.
Log. Methods Comput. Sci., 3(1).



Pesic, M., Schonenberg, H., and van der Aalst, W. M. P. (2007).
DECLARE: full support for loosely-structured processes.
In *Proc. of EDOC '07*.

References III



-  Rozinat, A. and van der Aalst, W. M. P. (2008).
Conformance checking of processes based on monitoring real behavior.
Inf. Syst., 33(1):64–95.
-  Torralba, Á., López, C. L., and Borrajo, D. (2016).
Abstraction heuristics for symbolic bidirectional search.
In *Proc. of IJCAI '16*.
-  van der Aalst, W. M. P. (2016).
Process Mining - Data Science in Action, Second Edition.
Springer.
-  Zhu, S., Tabajara, L. M., Li, J., Pu, G., and Vardi, M. Y. (2017).
Symbolic ltlf synthesis.
In *Proc. of IJCAI '17*.